

James Inah

612-489-9763 | jamesinah34@gmail.com | [linkedin.com/in/inah-james](https://www.linkedin.com/in/inah-james) | github.com/Troppy2 | St. Paul, MN

EDUCATION

University of Minnesota <i>B.S. Computer Science, College of Science and Engineering</i>	Enrolling Fall 2026 <i>Minneapolis, MN</i>
Century College <i>Transfer Coursework in Computer Science, GPA: 3.84/4.0</i> Coursework: Discrete Structures, Database Management Systems, OOP, C++, Python	Aug. 2025 – Present <i>White Bear Lake, MN</i>
Carnegie Mellon University (Open Learning) <i>Introduction to Modern AI (10-202) — Prof. Zico Kolter</i> Implemented a minimal LLM from scratch: custom tokenization, self-attention, and fine-tuning in PyTorch	Jan. 2026 – Present <i>Remote</i>

PROJECTS

Park&Go <i>Python, FastAPI, React, TypeScript, PostgreSQL, MapLibre GL, Docker</i> - Sustained 18.7 req/sec under 50 concurrent users with 0 failures across 1,109 requests and 150ms median latency, verified via Locust load testing; built a Redis caching layer with 5-min TTL across all read-heavy endpoints, dropping cached response times to sub-50ms in production - Engineered a 6-factor weighted scoring algorithm (cost, travel time, preferences, campus proximity, verification status, event adjacency) returning ranked parking recommendations in under 300ms with no cache; integrated Haversine distance computation and multi-mode travel time estimation (drive/walk) across 3 UMN campuses - Built 13 RESTful API endpoints across 6 routers with async FastAPI and SQLAlchemy AsyncSession — zero blocking DB calls; added Google OAuth SSO with JWT session management and rate limiting on submission endpoints via SlowAPI - Automated iCal ingestion from UMN's live campus calendar into PostgreSQL via APScheduler, enabling event-aware recommendations; instrumented Prometheus metrics on all routes plus 4 custom business counters for production observability - Shipped 123 passing tests (54 Python pytest + 69 frontend Vitest) spanning unit, integration, and E2E suites, covering geolocation mocks, nav store state transitions, and ETA population across a 10,000+ line codebase	Dec. 2025 – Present
Text Summarization, Seq2Seq Model <i>Python, PyTorch, NLTK, NumPy, scikit-learn</i> - Built a Seq2Seq model with Luong attention: 3-layer stacked LSTM encoder + 1-layer LSTM decoder trained on 500k+ Amazon food reviews - Achieved 90% summarization accuracy; condenses full reviews to single sentences with minimal loss - Custom preprocessing pipeline: HTML stripping, contraction expansion, Lancaster stemming, mode-based length selection, and from-scratch word2idx vectorization	Mar. 2026

TECHNICAL SKILLS

Languages: Python, Java, JavaScript/TypeScript, SQL, C++
Frameworks & Tools: FastAPI, React, PyTorch, SQLAlchemy, Docker, MapLibre GL, Tailwind CSS, Git
Databases: PostgreSQL, MySQL
Certification: Supervised Machine Learning: Regression & Classification (DeepLearning.AI)

STUDENT INVOLVEMENT

ColorStack <i>Member</i> - Engaged with a national community of 10,000+ Black and Latinx CS students focused on increasing representation in tech - Collaborated with peers on resume reviews, interview prep, and internship referral pipelines	Jan. 2026 – Present <i>Remote</i>
--	--------------------------------------